🦆 MomaGrid Cookbook Run

π Day — March 14

Decentralized AI Inference Grid — 3-GPU LAN Test

22/22	801s	13.4 min	http://192.168.0.177:9000
Recipes passed	Total elapsed	Wall time	Hub

Executive Summary

MomaGrid — a decentralized, open-source AI inference grid, running entirely on your own hardware, coordinating **3 GPU nodes**, serving **14 LLM models**, verified by cryptographic identity, with **22/22 cookbook recipes** all passing.

Built from scratch. Spec-driven. Pure Go. No cloud dependency. No API bills.

Run date: **2026-03-14 19:50:18** | Log: `/home/wengong/projects/digital-duck/momahub.go/cookbook/out/RUN_ALL_LOGGING-20260314-195018.md`

Recipe Results

ID	RECIPE	STATUS	ELAPSED	DESCRIPTION	NOTES
01	Hello SPL	✅	4.0s	Basic SPL query via grid	—
02	Multi-CTE Parallel	✅	26.1s	Parallel multi-CTE queries	—
03	Batch Translate	✅	2.3s	Batch translation — 4 languages	4/4 langs

ID	RECIPE	STATUS	ELAPSED	DESCRIPTION	NOTES
04	Benchmark Models	✅	8.3s	Model benchmark — latency & TPS	avg 22.2 tok/s · 3 models
05	RAG on Grid	✅	4.9s	RAG retrieval-augmented generation	142 tok
07	Stress Test	✅	56.4s	Stress test — sustained load	avg 30.8 tok/s · 15/15 tasks
08	Model Arena	✅	12.7s	Model arena — head-to-head	avg 34.3 tok/s · 3 models
09	Doc Pipeline	✅	8.3s	Document processing pipeline	284 tok
10	Chain Relay	✅	42.4s	Chain relay — sequential prompts	3 steps · 1707 tok
12	Tier-Aware Dispatch	✅	8.4s	Tier-aware dispatch by VRAM	avg 19.6 tok/s · 4/4 tasks
13	Multi-Agent Throughput	✅	27.6s	Multi-agent throughput benchmark	avg 12.8 tok/s · 31/60 tasks
15	Agent Failover	✅	22.7s	Agent failover & retry	20/20 tasks
16	Math Olympiad	✅	39.3s	Math olympiad — reasoning accuracy	qwen2-math=4/5, mathstral=4/5
17	Code Review Pipeline	✅	87.1s	Code review pipeline (3 stages)	3/3 steps
18	Smart Router	✅	48.6s	Smart router — query classification	avg 24.5 tok/s · 6/6 tasks
19	Privacy Chunk Demo	✅	15.5s	Privacy chunk — distributed doc analysis	3/3 chunks · assembled ✓
20	Overnight Batch	✅	18.4s	Overnight batch processing	avg 127.0 tok/s · 9/10 tasks
21	Language Accessibility	✅	8.4s	Language accessibility — 10 languages	9/10

ID	RECIPE	STATUS	ELAPSED	DESCRIPTION	NOTES
22	Rewards Report	✅	0.3s	Rewards report — operator credits	1077 tasks rewarded · 146.5 credits
23	Wake/Sleep Resilience	✅	60.6s	Wake/sleep resilience	20/20 tasks · 4 events
24	SPL Compiler Pipeline	✅	17.9s	SPL compiler pipeline	5 steps · 482 tok
25	Model Diversity	✅	281.2s	Model diversity — all models probed	14 models probed

Throughput Summary

RECIPE	AVG TOK/S
[04] Benchmark Models	**22.2**
[07] Stress Test	**30.8**
[08] Model Arena	**34.3**
[13] Multi-Agent Throughput	**12.8**
[18] Smart Router	**24.5**

Math Olympiad Scores

Difficulty: `medium`

MODEL	SCORE
qwen2-math	4/5
mathstral	4/5

Timing Breakdown

Top 5 slowest

#	RECIPE	ELAPSED
1	[25] Model Diversity	281.2s
2	[17] Code Review Pipeline	87.1s
3	[23] Wake/Sleep Resilience	60.6s
4	[07] Stress Test	56.4s
5	[18] Smart Router	48.6s

Top 5 fastest

#	RECIPE	ELAPSED
1	[22] Rewards Report	0.3s
2	[03] Batch Translate	2.3s
3	[01] Hello SPL	4.0s
4	[05] RAG on Grid	4.9s
5	[09] Doc Pipeline	8.3s

Generated 2026-03-14 20:04:55 by `cookbook/analyze_run_all.py` · MomaGrid — open-source decentralized AI inference